

# Persistent Memory Objects on the Cheap

International Conference on Supercomputing 2025.  
Salt Lake City, Utah, United States. June 11.

**Derrick Greenspan**, Naveed Ul Mustafa,  
Jongouk Choi, Mark Heinrich, Yan Solihin

CompArch & ARPERS research groups  
Cyber Security and Privacy Research Cluster



**College of Engineering  
and Computer Science**

UNIVERSITY OF CENTRAL FLORIDA




# Overview

- ➊ Introduction and Background
- ➋ Design
  - Light PMO (LPMO) Design
- ➌ Evaluation
  - LPMO Performance
  - CXL Performance
- ➍ Conclusion

# Persistent Memory (PM)


Example: Intel Optane, CXL Memory-Semantic SSDs



## Characteristics

# Persistent Memory (PM)

Example: Intel Optane, CXL Memory-Semantic SSDs



## Characteristics

- Slow write performance/Decent read performance

# Persistent Memory (PM)

Example: Intel Optane, CXL Memory-Semantic SSDs



## Characteristics

- Slow write performance/Decent read performance
- Poor multithreaded performance

# Persistent Memory (PM)

Example: Intel Optane, CXL Memory-Semantic SSDs



## Characteristics

- Slow write performance/Decent read performance
- Poor multithreaded performance ([more on this later](#))

# Persistent Memory (PM)

Example: Intel Optane, CXL Memory-Semantic SSDs



## Characteristics

- Slow write performance/Decent read performance
- Poor multithreaded performance ([more on this later](#))

Note: Devices with PM usually have DRAM as well



# Persistent Memory Objects (PMOs)

Primitives: pcreate()



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach()



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach() detach()



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach() detach() psync()



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach() detach() psync()

## Properties

- File-less
- Potentially pointer-rich
- Accessed via load-store instructions
- Metadata managed by kernel



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach() detach() psync()

## Properties

- File-less
- Potentially pointer-rich
- Accessed via load-store instructions
- Metadata managed by kernel

## Features



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach() detach() psync()

## Properties

- File-less
- Potentially pointer-rich
- Accessed via load-store instructions
- Metadata managed by kernel

## Features

- Fast with minimal metadata



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach() detach() psync()

## Properties

- File-less
- Potentially pointer-rich
- Accessed via load-store instructions
- Metadata managed by kernel

## Features

- Fast with minimal metadata
- Crash-consistency



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach() detach() psync()

## Properties

- File-less
- Potentially pointer-rich
- Accessed via load-store instructions
- Metadata managed by kernel

## Features

- Fast with minimal metadata
- Crash-consistency
- Security at rest



# Persistent Memory Objects (PMOs)

Primitives: pcreate() attach() detach() psync()

## Properties

- File-less
- Potentially pointer-rich
- Accessed via load-store instructions
- Metadata managed by kernel

## Features

- Fast with minimal metadata
- Crash-consistency
- Security at rest
- Integrity verification at rest




# Thread Scaling

Prior work exhibited poor thread scaling


# Thread Scaling

Prior work exhibited poor thread scaling



# Thread Scaling

Prior work exhibited poor thread scaling




## Reasons

- PMOs are hosted entirely in PM

# Thread Scaling

Prior work exhibited poor thread scaling



## Reasons

- PMOs are hosted entirely in PM
- Encryption and integrity verification on the critical path



# CXL Devices

CXL: SOTA for PM



CXL: SOTA for PM

## Compute Express Link

- Utilizes PCIe interface
- Direct access from CPU to memory
- Heterogeneous memory pools
- Can use PM or Volatile Memory



CXL: SOTA for PM

## Compute Express Link

- Utilizes PCIe interface
- Direct access from CPU to memory
- Heterogeneous memory pools
- Can use PM or Volatile Memory

## *Additional latency*

- From controller
- Comparable to NUMA



CXL: SOTA for PM

## Compute Express Link

- Utilizes PCIe interface
- Direct access from CPU to memory
- Heterogeneous memory pools
- Can use PM or Volatile Memory

## *Additional latency*

- From controller
- Comparable to NUMA

**Goal: High-Performance PMOs**



# CXL at-rest encryption

There is a way to do this...



# CXL at-rest encryption

There is a way to do this...

## CXL 3.1 Specification

- Trusted Execution Environments (TEE) Security Protocol (TSP)



# CXL at-rest encryption

There is a way to do this...

## CXL 3.1 Specification

- Trusted Execution Environments (TEE) Security Protocol (TSP)
- Range-based memory encryption



# CXL at-rest encryption

There is a way to do this...

## CXL 3.1 Specification

- Trusted Execution Environments (TEE) Security Protocol (TSP)
- Range-based memory encryption

I.e., **Transparent hardware encryption**



# CXL at-rest encryption

There is a way to do this...

## CXL 3.1 Specification

- Trusted Execution Environments (TEE) Security Protocol (TSP)
- Range-based memory encryption

I.e., **Transparent hardware encryption**

...more on this later.



# Overview

## ① Introduction and Background

## ② Design

- Light PMO (LPMO) Design

## ③ Evaluation

- LPMO Performance
- CXL Performance

## ④ Conclusion



# Threat Model

Goal: Protect at-rest data from disclosure/corruption



# Threat Model

Goal: Protect at-rest data from disclosure/corruption

## Out of Scope

- Side-channel attacks
- Data-remanence attacks (DRAM)



# Avoiding PM Pathologies

Prior work: PMO entirely in PM



# Avoiding PM Pathologies

Prior work: PMO entirely in PM

- Crash consistency simple



# Avoiding PM Pathologies

Prior work: PMO entirely in PM

- Crash consistency simple
- High latency, low write bandwidth



# Avoiding PM Pathologies

Prior work: PMO entirely in PM

- Crash consistency simple
- High latency, low write bandwidth

LPMO can exploit DRAM as cache **without** hardware support



# Avoiding PM Pathologies

Prior work: PMO entirely in PM

- Crash consistency simple
- High latency, low write bandwidth

LPMO can exploit DRAM as cache **without** hardware support  
DRAM as cache = Reconfigurable Memory



# Reconfigurable Memory Hierarchy

## Challenges



# Reconfigurable Memory Hierarchy

## Challenges

Which data should be placed in DRAM?



# Reconfigurable Memory Hierarchy

## Challenges

Which data should be placed in DRAM?

- All data in DRAM



# Reconfigurable Memory Hierarchy

## Challenges

Which data should be placed in DRAM?

- All data in DRAM
- Shadow in DRAM, primary in PM

Psync: Temporary Shadow Page (TSC) in PM, copy to Primary



# Reconfigurable Memory Hierarchy

## Challenges

Which data should be placed in DRAM?

- All data in DRAM
- Shadow in DRAM, primary in PM

Psync: Temporary Shadow Page (TSC) in PM, copy to Primary



# Reconfigurable Memory Hierarchy

## Challenges

Which data should be placed in DRAM?

- All data in DRAM
- Shadow in DRAM, primary in PM

Psync: Temporary Shadow Page (TSC) in PM, copy to Primary

What data should be encrypted?

## Challenges

Which data should be placed in DRAM?

- All data in DRAM
- Shadow in DRAM, primary in PM

Psync: Temporary Shadow Page (TSC) in PM, copy to Primary

What data should be encrypted?

- Primary and shadow page encrypted



# Reconfigurable Memory Hierarchy

## Challenges

Which data should be placed in DRAM?

- All data in DRAM
- Shadow in DRAM, primary in PM

Psync: Temporary Shadow Page (TSC) in PM, copy to Primary

What data should be encrypted?

- Primary and shadow page encrypted
- Primary and shadow in plaintext

## Challenges

Which data should be placed in DRAM?

- All data in DRAM
- Shadow in DRAM, primary in PM

Psync: Temporary Shadow Page (TSC) in PM, copy to Primary

What data should be encrypted?

- Primary and shadow page encrypted
- Primary and shadow in plaintext
- Shadow in plaintext, primary in ciphertext




# Reconfigurable Memory Hierarchy (Part 2)

Optane PM uses local DRAM

# Reconfigurable Memory Hierarchy (Part 2)

Optane PM uses local DRAM



CXL can place system on either side of memory expander



# Reduce page fault latency

Prior work: Demand Faulting

Why not **predict** when pages are needed?



# Reduce page fault latency

Prior work: Demand Faulting

Why not **predict** when pages are needed?

Example Solution: Stream Buffer

- On fault, predict next  $X$  sequential pages (**depth**)
- Works well for access patterns amenable to prediction

# Overview

## ① Introduction and Background

## ② Design

- Light PMO (LPMO) Design

## ③ Evaluation

- LPMO Performance
- CXL Performance

## ④ Conclusion

# Evaluation

## Evaluated Benchmarks

- Microbenchmarks
  - 2d Convolution (2dConv)
  - Gaussian Elimination (Gauss)
  - LU Decomposition (LU)
  - Tiled Matrix Matrix Multiplication (TMM)
- Filebench (Fileserver, VarMail, WebProxy, WebServer)
- LMDB

# Evaluation


## Evaluated Benchmarks

- Microbenchmarks
  - 2d Convolution (2dConv)
  - Gaussian Elimination (Gauss)
  - LU Decomposition (LU)
  - Tiled Matrix Matrix Multiplication (TMM)
- Filebench (Fileserver, VarMail, WebProxy, WebServer)
- LMDB


Component	Specifications
MB	Supermicro X11DPi-NT
CPU	2×Intel Xeon Gold 6230 (20 cores)
DRAM	4 × 32GiB DDR4 @ 2666MHz
PM	4 × 128GiB Intel Optane DIMM
OS	AlmaLinux 9.0; Linux 5.15.157

# LPMO Performance


△ O, ◎ D, □ D2, × D4, ▷ D8



(a) Gauss




(b) LU



(c) TMM


All benchmarks have better thread scaling!

# LPMO Performance




- DRAM reduces execution time by  $\approx 21\%$

# LPMO Performance




- DRAM reduces execution time by  $\approx 21\%$
- IV + Prediction faster than original GPMO design w/o IV

# LPMO Performance - Filebench




- Only 1.19 $\times$  faster with DRAM

# LPMO Performance - Filebench



- Only 1.19 $\times$  faster with DRAM
- 1.81 $\times$  **faster with page prediction**


# LPMO Performance - Filebench



- Only  $1.19\times$  faster with DRAM
- $1.81\times$  faster with page prediction**
- $1.37\times$  faster with page prediction & IV


# CXL Performance

Perform same tests, but emulate CXL latency




- Use opposite-node Optane
- Near configuration: cache allocated from local node
- Far configuration: cache allocated from opposite node

# CXL Performance



- With CXL alone: 50% slower than original
- With DRAM: 20% faster (despite CXL latency)

# CXL Performance - Filebench





# Overview

## ① Introduction and Background

## ② Design

- Light PMO (LPMO) Design

## ③ Evaluation

- LPMO Performance
- CXL Performance

## ④ Conclusion



# Conclusion

## LPMO

- Software-based DRAM caching
  - Up to  $1.25\times$  faster



# Conclusion

## LPMO

- Software-based DRAM caching
  - Up to  $1.25\times$  faster
- Predictive Decryption
  - Up to  $1.81\times$  faster



# Conclusion

## LPMO

- Software-based DRAM caching
  - Up to  $1.25\times$  faster
- Predictive Decryption
  - Up to  $1.81\times$  faster

## CXL

- Introduced Reconfigurable Memory Hierarchy
- CXL latency can be masked by LPMO optimizations



## Q & A

Thank You!

Any questions?